

# Setting up Kubernetes from VMs

Lets setup a basic Kubernetes cluster using Virtual Machines (VMs).

## Plan

Here's the overall plan:

- ✓ Install VirtualBox
- ✓ Install AlmaLinux 9.6 Minimal ISO (choose ARM64 for Apple Silicon)
- ✓ Spin up two VMs: one for the master node, and one for the worker node.
- ✓ Install Kubernetes using `kubeadm` .

## Workflow

Let's break down the steps:

- ✓ **Install VirtualBox**
- ✓ **Download ISO**
  - Download the AlmaLinux 9.6 Minimal ISO.
- ✓ **Configure VirtualBox Network Settings**
  - Chosen **NAT Network** as the primary network mode.
  - Created a network named `k8s-network` .
  - Defined the CIDR (Classless Inter-Domain Routing): `10.0.2.0/24` .
    - `/24` means that 24 out of the 32 bits are used for the network address, leaving 8 bits for host addresses.
    - This gives us the IP address range: `10.0.2.0 - 10.0.2.255` .
  - Disabled DHCP to assign static IP addresses to our VMs.
- ✓ **Create VM1: Master Node** (4GB RAM, 15GB storage)
  - Name: `k8s-master1`
  - ISO: Path to the downloaded ISO.
  - Type: Linux
  - SubType: Oracle Linux
  - Version: Oracle Linux 9.x
  - Network: Attached to **NAT Network** ( `k8s-network` )
  - **Installing OS:**
    - Created a root user password.
    - Created a new sudo user `k8s` with a password.
    - Disabled KDUMP.

- Configured a static IP address:
  - IP: 10.0.2.10
  - NetMask: 255.255.255.0
  - Gateway: 10.0.2.1
  - DNS: 8.8.8.8
- Began the installation.
- Rebooted the system.
- Installation completed, got the shell.
- Checked IP address: `ip addr show` (should show 10.0.2.10/24 ).
- Checked internet connectivity: `ping hrushispace.com` .
- Checked OS info: `cat /etc/os-release` .
- Checked memory: `free -h` .
- Checked disk: `df -h` .
- Checked CPU: `lscpu` .

#### ☒ **Create VM2: Worker Node** (4GB RAM, 15GB storage)

- Name: k8s-worker1
- ISO: Path to the downloaded ISO.
- Type: Linux
- SubType: Oracle Linux
- Version: Oracle Linux 9.x
- Network: Attached to **NAT Network** ( k8s-network )
- **Installing OS:**
  - Created a root user password.
  - Created a new sudo user k8s with a password.
  - Disabled KDUMP.
  - Configured a static IP address:
    - IP: 10.0.2.11
    - NetMask: 255.255.255.0
    - Gateway: 10.0.2.1
    - DNS: 8.8.8.8
  - Began the installation.
  - Rebooted the system.
  - Installation completed, got the shell.
  - Checked IP address: `ip addr show` (should show 10.0.2.11/24 ).
  - Checked internet connectivity: `ping hrushispace.com` .
  - Checked OS info: `cat /etc/os-release` .
  - Checked memory: `free -h` .
  - Checked disk: `df -h` .

- Checked CPU: `lscpu` .

### ✅ Verify VM Communication

- On `k8s-master1` : `ping 10.0.2.11`
- On `k8s-worker1` : `ping 10.0.2.10`

### ✅ Check Active Services

- On each VM: `systemctl list-units --type=service --state=actives`

### ✅ Disable Firewall

- By default, the firewall is active, which can block necessary communication. While you can open specific ports, disabling it for this basic setup is simpler.
- On both master and worker nodes:

```
sudo systemctl disable firewalld --now
```

### ✅ Disable Swap

- Swap is virtual memory on disk.
- Kubernetes prefers real memory for predictable resource allocation.
- Disable swap:

```
sudo swapoff -a
```

- To prevent swap from being enabled on reboot, comment it out in `/etc/fstab` :

```
sudo sed -i '/ swap / s/^(.*)$/#\1/g' /etc/fstab  
free -h
```

### ✅ Adding Kernel Modules

- Think of kernel modules as extensions for your Linux OS, providing extra functionality when needed. Kubernetes networking relies on `overlay` and `br_netfilter` .
- Go to `/etc/modules-load.d/` .
- Create a file `k8s.conf` and add these two modules:

```
overlay  
br_netfilter
```

- `overlay` : Used by container runtimes like Podman for efficient management of container image layers.
- `br_netfilter` : Crucial for allowing network traffic through Linux bridges to be processed by the firewall (iptables), which Kubernetes often uses for its networking.
- Load the modules using `modprobe` :

```
sudo modprobe br_netfilter
sudo modprobe overlay
```

- Alternatively, rebooting the VMs will automatically load these modules.

### ✓ **Configure sysctl for Network**

- For communication between pods and to enable firewall rules for pod traffic, configure `sysctl`.
- Create `k8s.conf` inside `/etc/sysctl.d/` and add these lines:

```
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
```

- `net.bridge.bridge-nf-call-iptables = 1` : Enables iptables rules for traffic traversing Linux bridges.
- `net.bridge.bridge-nf-call-ip6tables = 1` : Same as above, but for IPv6.
- `net.ipv4.ip_forward = 1` : Allows IP forwarding, necessary for inter-pod communication across different nodes.

- Reload the `sysctl` configuration:

```
sudo sysctl --system
```

### ✓ **Install CRI-O Container Runtime**

- Create a repository file for CRI-O:

```
cat <<EOF | sudo tee /etc/yum.repos.d/cri-o.repo
[cri-o]
name=CRI-O
baseurl=[https://download.opensuse.org/repositories/isv:/cri-o:/stable:/$(http:
enabled=1
gpgcheck=1
gpgkey=[https://download.opensuse.org/repositories/isv:/cri-o:/stable:/$(https
EOF
```

- Install CRI-O and `container-selinux` :

```
sudo dnf install -y container-selinux cri-o
sudo systemctl start crio.service
```

- CRI (Container Runtime Interface) is an abstraction layer that allows Kubernetes to work with different container runtimes.

## ✓ Install Kubernetes

- Add the Kubernetes repository:

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=[https://pkgs.k8s.io/core:/stable:/v1.29/rpm/](https://pkgs.k8s.io/core
repo_gpgcheck=1
gpgcheck=1
gpgkey=[https://pkgs.k8s.io/core:/stable:/v1.29/rpm/repodata/repomd.xml.key](ht
enable=1
EOF
```

- Install `kubelet`, `kubeadm`, and `kubectl`:

```
sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
```

## ✓ Configure Master Node

- Initialize the Kubernetes control plane on the master node:

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --cri-socket unix:///var/run,
```

- Configure `kubectl` to work with your cluster:

```
mkdir -p /home/k8s/.kube
sudo cp /etc/kubernetes/admin.config /home/k8s/.kube/config
sudo chown k8s:k8s /home/k8s/.kube/config
```

## ✓ Configure Worker Node

- You'll see a `kubeadm join` command output by the `kubeadm init` command on the master node. Copy and run that command on the worker node. It will look something like this:

```
kubeadm join <your_master_ip>:<port> --token <token> --discovery-token-ca-cert-l
```

## ✓ Install a CNI Plugin

- Kubernetes needs a Container Network Interface (CNI) plugin to enable communication between pods. We'll install Calico:

```
kubectl apply -f [https://raw.githubusercontent.com/projectcalico/calico/v3.27.0
```

- Once the CNI plugin is deployed, your worker node should be able to join the cluster.

```
[k8s@10 kubernetes]$ kubectl get nodes
NAME           STATUS    ROLES          AGE    VERSION
10.0.2.10      Ready    control-plane   66m    v1.29.15
10.0.2.11      Ready    <none>         28m    v1.29.15
[k8s@10 kubernetes]$
```

## VirtualBox Network Modes

Understanding VirtualBox network modes is key:

- **1. NAT (Default)** - [ Room with no door between rooms, only enter/exit from/to outside ]
  - VM gets internet through the host.
  - VMs are isolated from each other.
  - VM can access the internet.
  - VMs cannot reach each other.
- **2. NAT Network** - [ Rooms connected by Hallway with access to outside ]
  - VMs share a virtual network and can communicate.
  - VM can access the internet.
  - Host cannot directly reach the VM.
  - VMs can reach each other.
- **3. Host-only Adapter** - [ Same as NAT but no access to outside ]
  - Creates a private network between the host and VMs.
  - VM cannot access the internet.
  - Host can reach the VM.
  - VMs can reach each other.
- **4. Bridged Adapter** - [ Each room has access to outside ]
  - VM connects directly to your physical network.
  - VM can access the internet.
  - Host can reach the VM.
  - VMs can reach each other.

## DHCP

**Dynamic Host Configuration Protocol** - An automatic IP address assignment service that gives devices network settings when they connect.

## Modprobe

**Modprobe** is a command-line utility in Linux used to add and remove modules from the Linux kernel.